

Self-supervised pretraining for low resource languages in Speech processing

Julien Hauret & Nicolas Dufour

Algorithms for speech and natural
language processing

MVA 2020-2021

ENS Paris-Saclay


March 25th, 2021

I. Presentation of the task

Phoneme recognition



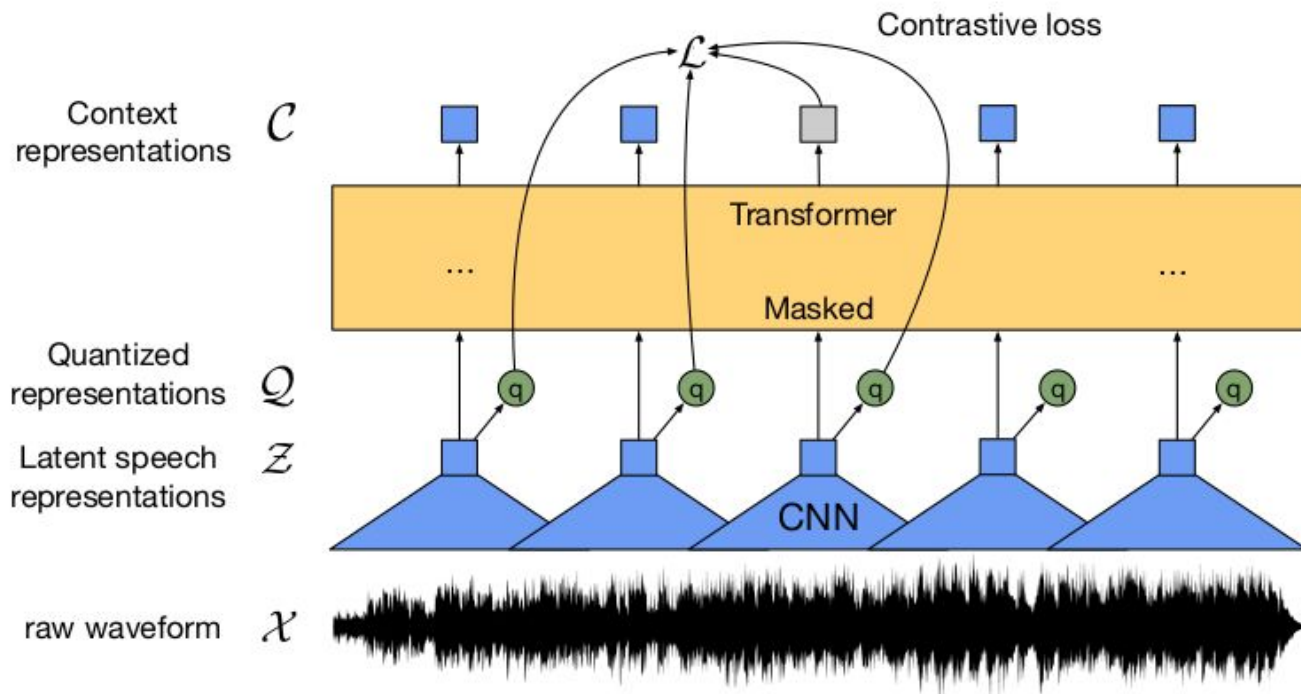
ʒ a d ɔ ʁ l ə k a m ã b ɛ ʁ t



II. Models And metrics

January 22th, 2021

Wav2Vec2 2.0 architecture




CTC Loss


For an input,
like speech



Predict a
sequence of
tokens



Merge repeats,
drop ε




Final output

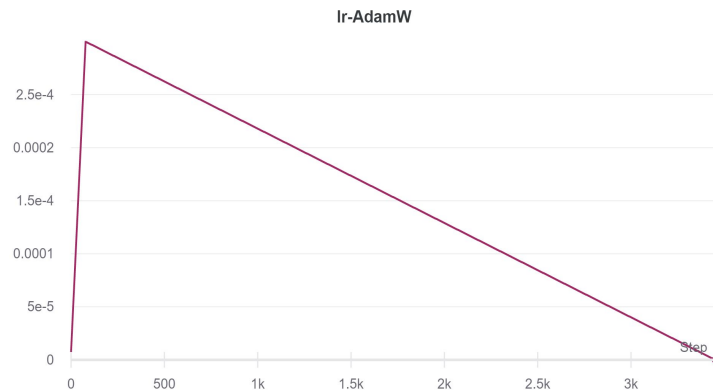


Picture take from <https://distill.pub/2017/ctc/>

Data preprocessing

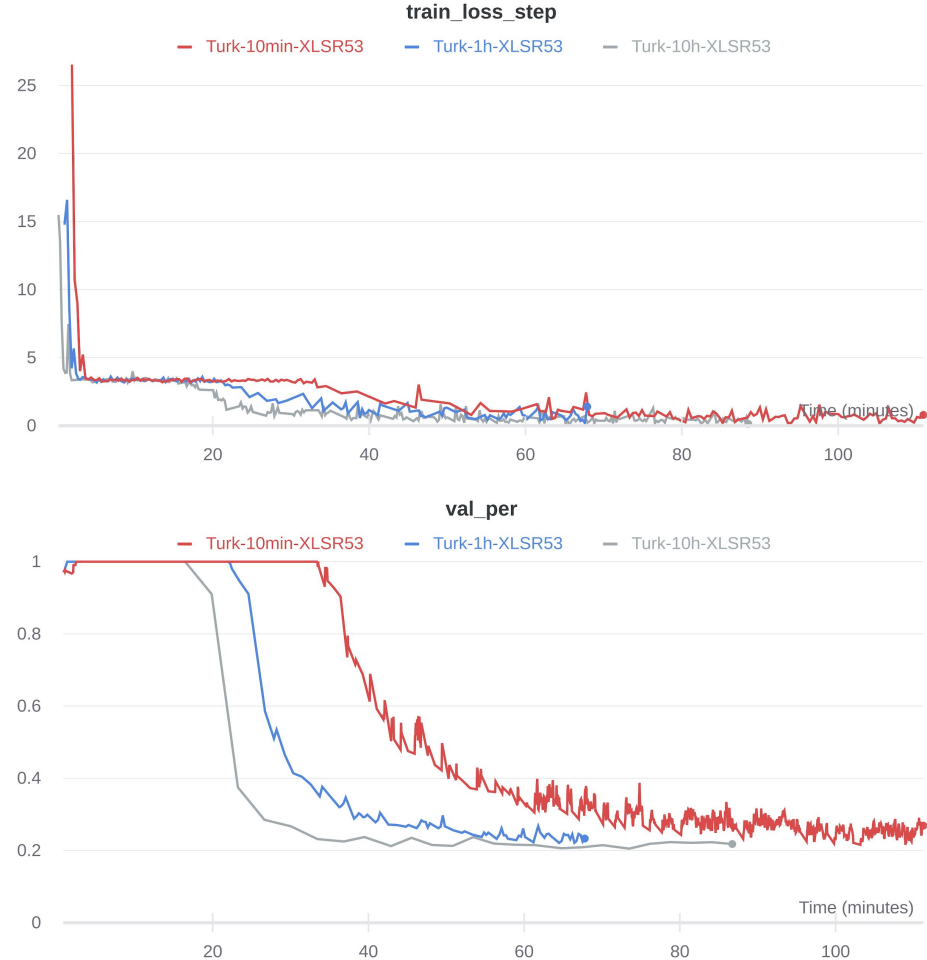


Training procedure



- AdamW with max lr of $3e-4$
- Model selection on val PER
- 30 epochs

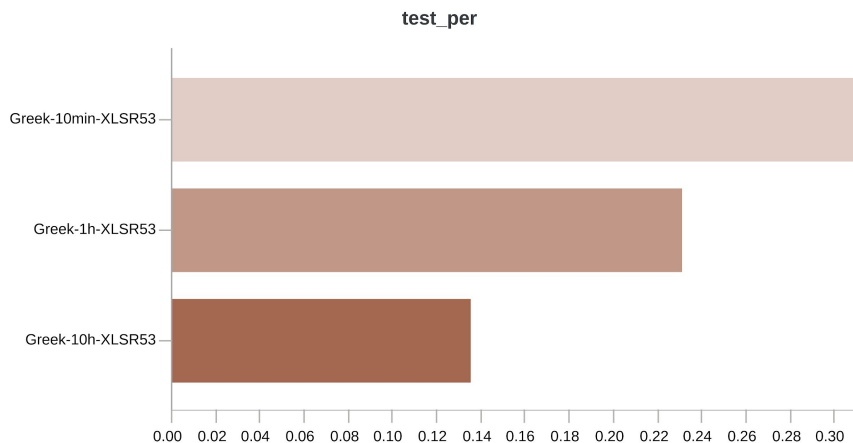
Training curves :




III.Experiments

January 22th, 2021

Impact of data quantity : low proportion in XLSR-53

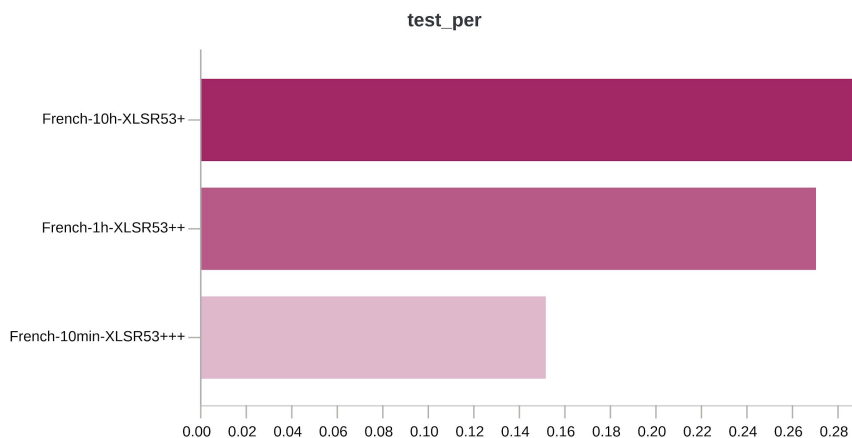


Greek




Turk

Impact of data quantity : high proportion in XLSR-53




French




Spanish

What's the best pretrained model?




Czech training

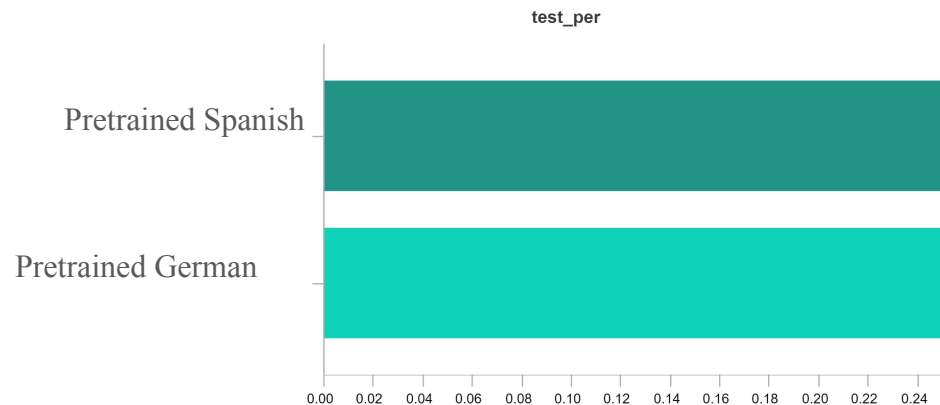


Greek training

Does it matter on which language we pretrain?



Portuguese training



Dutch training

Conclusion

Conclusion

We've train multiple models and studied the properties of fine-tuning for low ressources languages and draw the following conclusion:

- When the language has a small contribution on XLSR-53, training the more data the better.
- When the language has a bigger contribution on XLSR-53, we didn't observe the same correlation.
- We proved that pretraining is without a doubt useful but didn't manage to prove that XLSR-53 was more adapted to fine-tuning than a model only trained on English
- Finally, pre-training on close languages seems to have a positive effect on fine-tuning

Thank you for listening

March 25th, 2021

Julien Hauret & Nicolas Dufour